

```
In [51]: print("Author: Louis-Antoine Etchian")
```

Author: Louis-Antoine Etchian

```
In [52]: %matplotlib inline

import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
```

Plotting with pandas

We use the standard convention for referencing the matplotlib API ... We provide the basics in pandas to easily create decent looking plots.

https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html
(https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html).


Let's load the heart attack dataset

```
In [55]: data_names=np.array(["mpg", "cylinders", "displacement", "horsepower", "weight",  
data=pd.read_csv('auto-mpg.data', delim_whitespace=True, header=None, names = data_names)
```

Plotting all columns, works, but does not provide a lot of insight.


```
In [56]: data.plot()
```

Out[56]: <AxesSubplot:>



Let's look at the age distribution (a histogram)

```
In [57]: data['mpg'].plot.hist();
```



How many male and female samples do we have?

```
In [58]: data.origin.value_counts()
```


```
Out[58]: 1    249
          3     79
          2     70
          Name: origin, dtype: int64
```

Notice that we accessed the gender column with dot notation. This can be done whenever the column name is 'nice' enough to be a python variable name.

Do we have similar ages in females and males?


Plotting two histograms for each gender side beside directly from the dataframe:

```
In [59]: #Same focus than the previous documents (mpg by origin)
          axs = data.hist(column='mpg', by='origin')
```




To format this plot, we can work on the axes (array) that is returned by the plot call. We use Matplotlib object oriented interface methods to do this

```
In [62]: axs = data.hist(column='mpg', by='origin')
#axs[0].set(title='Origin #1', ylim=[1, 40])
#axs[1].set(title='Origin #2', ylim=[1, 40])
#axs[2].set(title='Origin #3', ylim=[1, 40])
```



Is age and blood pressure correlated? Maybe it is different for females and males?
Let's have a look with a scatter plot.

```
In [32]: data.plot.scatter('mpg', 'weight', c='origin', colormap='winter');
```




According to:

<https://stackoverflow.com/questions/43578976/pandas-missing-x-tick-labels>
(<https://stackoverflow.com/questions/43578976/pandas-missing-x-tick-labels>)

the missing x-labels are a pandas bug.

Workaround is to create axes prior to calling plot

```
In [35]: fig, ax = plt.subplots()
data.plot.scatter('mpg', 'weight', c='origin', colormap='viridis', ax=ax);
```




It is a bit annoying that there is a colorbar, we know gender is categorical.

One way to avoid the colorbar is to loop over the categories and assign colors based on the category.

See: <https://stackoverflow.com/questions/26139423/plot-different-color-for-different-categorical-levels-using-matplotlib> (<https://stackoverflow.com/questions/26139423/plot-different-color-for-different-categorical-levels-using-matplotlib>)

```
In [40]: colors = {1: 'tab:blue', 2: 'tab:orange', 3: 'tab:green'}
fig, ax = plt.subplots()
for key, group in data.groupby(by='origin'):
    group.plot.scatter('mpg', 'weight', c=colors[key], label=key, ax=ax);
```



Seaborn

Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics.


<http://seaborn.pydata.org/index.html> (<http://seaborn.pydata.org/index.html>)

Seaborn is usually imported as `sns`

```
In [41]: import seaborn as sns
```


Let's re-create the histograms by gender with seaborn with the figure level `displot()` function.

```
In [42]: # Use origin to split age into columns  
sns.displot(x='mpg', col='origin', data=data);
```




We can display the counts in the same plot, one on top of the other.

```
In [43]: # Use origin to color (hue) in the same plot  
sns.displot(x='mpg', hue='origin', data=data);
```




To have an idea of the split between male and female, we can stack the counts, adding up to total.

```
In [44]: sns.displot(x='mpg', hue='origin', data=data, multiple='stack');
```



We can look at the differences in ages with a boxplot too

```
In [45]: sns.boxplot(x='origin', y='mpg', data=data);
```




Let's re-create the scatter plot to see if age and blood pressure are correlated by gender.

To make the legend show strings we will create a gender string column with female and male strings rather than 0 and 1.


```
In [46]: data['origin_str'] = data['origin'].replace([1,2,3], ['1','2','3'])
```

```
In [47]: ax = sns.scatterplot(x='mpg', y='weight', data=data, hue='origin_str')
```



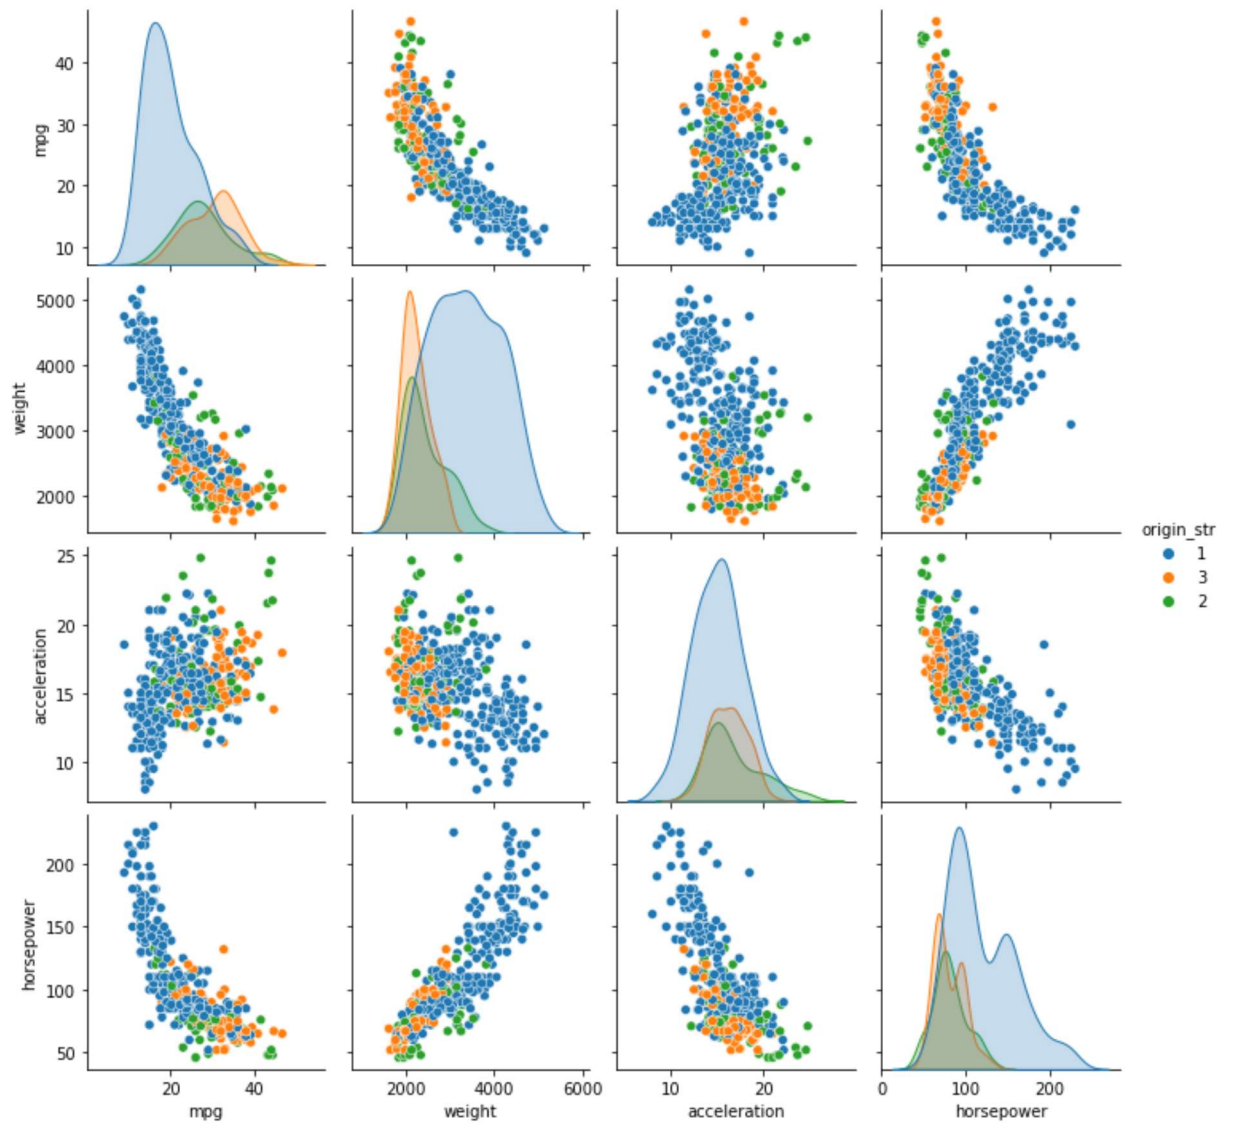
Adding a regression line helps with visualizing the relationship

```
In [48]: ax = sns.lmplot(x='mpg', y='weight', data=data, hue='origin_str')
```




Maybe there are other correlations in the data set. Pairplot is a great way to get an overview


```
In [49]: sns.pairplot(data, vars=['mpg', 'weight', 'acceleration', 'horsepower'], hue='origin_str')
```



As an alternative, we can visualize the correlation matrix as a heatmap

```
In [50]: g = sns.heatmap(data[['mpg', 'weight', 'acceleration', 'horsepower']].corr(method='spearmanr', annot=True))
```



There are nice tutorials on the Seaborn website, be sure to check these out.

```
In [ ]:
```